

CSCI 5521: Introduction to Machine Learning (Spring 2022)¹

Homework 3

Due date: Apr 6, 2022 11:59pm

1. **(30 points)** Consider the Multilayer Perceptron (MLP) for binary classification described in section 11.7.2 in the textbook. Let's look at the regularized version of MLP when the activation function of each hidden unit becomes the reLU function $\text{reLU}(x) = \max(0, x)$, and the activation function of the output unit is the sigmoid function. In the regularized version, the error function becomes the following:

$$E(W, v|X) = - \sum_{t=1}^N [r^t \log y^t + (1 - r^t) \log(1 - y^t)] + \sum_{h=1}^H \|w_h\|_2^2,$$

where $y^t = \text{sigmoid}(\sum_{h=1}^H v_h z_h^t + v_0)$ and $z_h^t = \text{reLU}(w_h^T x^t + w_{h0})$. Derive the update equations of the regularized MLP using the given activation functions. Note that you need to show the update equations for all trainable parameters.

Hint 1: Given $y = \text{sigmoid}(\alpha) = 1/(1 + e^{-\alpha})$, the derivative $\frac{\partial y}{\partial \alpha} = y(1 - y)$.

Hint 2: Given $\text{reLU}(f(x)) = \max(0, f(x))$, the derivative of $\text{reLU}(f(x))$ is given by $f'(x)$ if $f(x) > 0$, and 0 otherwise.

2. **(30 points)** Build a (Multilayer) Perceptron to recognize a certain area of the plane. That is, the Perceptron should output a “1” if the input vector lies in the shaded region.

- (a) Determine the vector of coefficients $w = [w_0, w_1, w_2]^T$ for a single layer perceptron of the form in Figure 1 to recognize the area in Figure 2 and again for Figure 3 shaded blue. Use a step-function as the non-linear activation function at designated nodes:

$$s(a) = \begin{cases} 1, & \text{if } a > 0 \\ 0, & \text{otherwise} \end{cases}$$

- (b) Determine coefficients $\mathbf{W} = \begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{20} & w_{21} & w_{22} \end{bmatrix}^T$, $v = [v_0, v_1, v_2]^T$ in the 2-layer Perceptron of the form in Figure 4 to recognize the shaded region in Figure 5.

Hint: The shaded region in Figure 5 equals to the intersection of the regions of Figures 2 & 3.

¹Instructor: Catherine Qi Zhao. TA: Prithvi Raj Botcha, Shi Chen, Suzie Hoops, James Yang, Yifeng Zhang. Email: csci5521.s2022@gmail.com




Figure 1: Single-layer perceptron. The non-linearity is a step function yielding a discrete value 1/0.




Figure 2: Accept shaded region.




Figure 3: Accept shaded region.

3. **(40 points)** In this programming exercise you will implement a Multilayer Perceptron (MLP) for optical-digit classification. You will train your MLP using the `optdigits_train.txt` data, tune the number of hidden units using the `optdigits_valid.txt` data, and test the prediction performance using the `optdigits_test.txt` data. For each file, the first 64 columns correspond to the features for different samples while the last one stores the labels. Features in each matrix should be normalized as $X_{norm} = \frac{(X - \mu_{trn})}{\sigma_{trn}}$. Notice that μ and σ , the mean and the standard deviation, are always calculated from the training set (even when normalizing the validation and test set).

- (a) Implement a MLP with 1 hidden layer for classifying the 10 digits (read section 11.7.3 in the textbook, note that each update should consider all training samples), use the tanh activation function $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$ for the hidden layer, and the softmax activation function $\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_k e^{x_k}}$ for the output layer. The error




Figure 4: Multi-layer perceptron.




Figure 5: Accept shaded region (intersection of Figures 2 & 3)

function is the cross-entropy loss:

$$E(r^t, y^t) = - \sum_i^C r_i^t \log y_i^t \quad (1)$$

where y^t is the predicted probabilities for different classes, C is the number of candidate labels. r^t is the one-hot vector for ground truth label, $r_i^t = 1$ if the current sample belongs to the i_{th} class and 0 otherwise.

Try MLPs with $H = 4, 16, 20, 24, 32$ and 48 hidden units. **Report the validation accuracy by the number of hidden units. How many hidden units should we use? Report the accuracy on the test set using this number of hidden units.**

Hint 1: Given $y_i = \text{softmax}(\alpha_i) = \frac{e^{\alpha_i}}{\sum_k e^{\alpha_k}}$, the derivative $\frac{\partial y_i}{\partial \alpha_j} = y_i(\delta_{ij} - y_j)$, where $\delta_{ij} = 1$ if $i = j$ and 0 otherwise.

Hint 2: Instead of manually checking the values of i and j , you could write down the whole update equation and simplify it to cancel out δ_{ij} .

- (b) Train two MLPs with 2 and 3 hidden units, respectively. Visualize the data by values of their hidden units. **Compare the 2-D and 3-D plots and explain the results in the report.** The code for plotting is included in `visualization.py`, and you do not need to modify the file.

We have provided the skeleton code `MyMLP.py` for implementing the algorithm. `MyMLP.py` is written in a *scikit-learn* convention, where you have a `fit` function for model train-

ing and a `predict` function for generating predictions on given samples. To verify your implementation, call the main function `hw3.py`.

Submission

- **Things to submit:**

1. `hw3_sol.pdf`: a document containing all your answers for the written questions (including answers and plots for problem 3).
2. `MyMLP.py`: a Python source file containing your implementation of the MLP algorithm (`class MLP`, `def tanh` and `def softmax`), and functions for data processing (`class Normalization` and `def process_label`). Use the skeleton file `MyMLP.py` found with the data on the class website, and fill in the missing parts. The `fit` function should take both training and validation samples as inputs, update the model parameters based on training samples and determine convergence based on validation accuracy. It should also return the best validation accuracy. The `predict` function should take features as inputs and return the predicted class labels. The `get_hidden` function should extract features computed at the hidden layers for given inputs. The `Normalization` class should compute the statistics of training samples with its `fit` function, and normalize given samples using the `normalize` function. The `process_label` function should convert class labels into one-hot vectors.

- **Submit:** All material must be submitted electronically via Gradescope. **Note that There are two entries for the assignment, *i.e.*, Hw3-Written (for `hw3_sol.pdf`) and Hw3-Programming (for a zipped file containing the Python code), please submit your files accordingly.** We will grade the assignment with vanilla Python, and code submitted as iPython notebooks will not be graded.